

Felice Antonio Merra, Ph.D.

Data & AI Leader | Senior Data Science and Analytics Manager

+39 340 42 13 560 · merrafelice@gmail.com · merrafelice.github.io · linkedin.com/in/feliceantoniomerra

EXECUTIVE SUMMARY

Data and AI leader with a Ph.D. in Computer Science and 10 years across applied AI, research and software engineering — five of them building and leading teams at Amazon and Cognism. Took an inherited science team from **zero models in production to nine in under two years**, and grew it from three people to a 12-person, two-function organisation spanning Data Science, Data Analytics and co-owned MLOps planning. The pattern across both companies is the same: build the measurement and delivery infrastructure that lets a stalled technical work ship — top-notch generative model in production, definition of evaluation frameworks, an RFC-to-production process, a hiring loop, and the data quality OKR definitions and programme execution. Operates at executive level — weekly with the VP of Data, monthly with the CPTO, core contributor to quarterly and multi-year investment planning — while remaining hands-on across LLMs, multi-agent systems and large-scale ranking and retrieval. 40+ publications in top-tier venues; Industry Chair, RecSys 2026. Programme Committee member at RecSys, SIGIR, and top-tier conferences.

LEADERSHIP & ORGANISATIONAL IMPACT

- **Turned research into delivery — 0 to 9 models in production in under two years.** Inherited a science team decoupled from engineering with nothing shipped; designed the RFC-to-production process from scratch: a three-month cycle with a mid-point checkpoint handover letting MLOps and ML Engineers build infrastructure in parallel rather than sequentially. 92–93% of deliveries on time.
- **Built and scaled the organisation from 3 to 12, in two stages.** Inherited a team of three; grew it to seven, then — following promotion to Senior Manager on delivery, leadership and scope — to a 12-person, two-function org. 3 hires, 4 promotions, plus a Principal Scientist attracted from another team: a level that would not normally report into this role.
- **Unified Data Science and Data Analytics under one function.** Proposed and won approval for the reorg, then installed a single prepared planning cadence in place of recurring status meetings. Data analytical tasks delivery-as-committed rose from an estimated 50–60% to a measured 95% within three months.
- **Built the science hiring bar from zero.** Replaced an informal process with a four-stage loop: phone screen, live coding with AI code assistants permitted, a two-part science interview (ML foundations + applied problem), and a joint behavioural panel. Hired from Amazon and Imperial College London.
- **Engineered growth paths rather than waiting for headcount.** Scoped a growth activity for a mid-level scientist that progressed from exploration to POC, to department project, to company flagship initiative; promoted them to Senior on delivered impact.
- **Executive influence.** Weekly with the VP of Data, monthly with the CPTO, core contributor to quarterly and multi-year investment planning, data strategy and buy-vs-build decisions; has presented to the executive team. DRI for multiple company “big rocks” strategic initiatives, including A-adoption.
- **DRI for the company's AI adoption roadmap**, including the evaluation tooling that measures the efficiency it delivers. Drove agentic tooling (Claude Code, OpenAI ADK) to full adoption across the data department, authored project-management and requirements-definition skills for the company's internal plugin, and onboarded managers in other functions.
- **Proposed an LLM-as-a-Judge annotation component delivering 100× the annotation volume of the previous human process at up to 1/1000 of the cost.** Redesigned the scientist workflow to be AI-first end to end — LLM-draft-then-review at each step — supported by knowledge-sharing sessions and a tagged Slack channel whose routines compile discussion into shared documentation.

EXPERIENCE

Senior Science and Analytics Manager | Cognism

February 2026 – Present · Remote (IT)

- Lead and scale the Science and Analytics organisation of an AI-first SaaS company — 12 people across Data Science and Data Analytics — owning people leadership, organisational design, hiring, technical strategy and delivery, with co-ownership of MLOps planning and roadmap.
- Partner with the CPTO on company-wide AI-strategy & adoption, long-term planning and investment priorities, translating emerging AI capabilities into scalable products and measurable business value.
- **Contributed to the company's first data quality OKR** and established precision and match rate as paired, separately tracked metrics and running education sessions across the business on the trade-off between them. Delivered up to ~30% improvement in precision across multiple projects keeping the coverage the same or higher.
- **Parameterised delivered models so product teams tune precision against match rate per use case** — the business sets the trade-off rather than inheriting a fixed one from science.
- **Shipped the company's first production agentic workflow**, now the de facto process for a company-wide programs: replaced the manual workload of two teams, which moved from executing to supervising, with reviewer agents built in to drive accuracy upward. Presented outcomes at company town hall.

Data Science Manager | Cognism

December 2024 – January 2026 · Remote (IT)

- Led and scaled the Data Science team from 3 to 7 scientists delivering AI/ML solutions for international SaaS initiatives — from technical strategy and solution design through deployment and production adoption.
- **Proposed and built Orion, Cognism's AI-powered data collection and validation engine** (publicly announced, Q1 2025) — a multi-agent system that collects, validates and confidence-scores data at scale, delivered in production.
- **Won and delivered a build-over-buy decision on data acquisition.** Recommended building rather than licensing to the VP of Data and the Principal PM owning Data Acquisition; delivered on time. Built as a parameterised framework, since reused for later build decisions — one doubling coverage of a key data asset.
- **Established scientific rigour where there was none.** Retired and rebuilt legacy models at higher accuracy and coverage; built a reusable evaluation package adopted across all science projects, covering accuracy metrics, statistical deviance and feature-parity tests. Set technical direction and delivery standards between Science and MLOps.
- Led projects spanning multi-agent systems, multilingual AI, information retrieval and advanced ML. Own data quality and metrics-layer strategy end-to-end on Amazon Athena, Iceberg and DynamoDB; delivered transformer-based quantised ranking and retrieval models serving tens of millions of requests per day on AWS (Triton).
- **Made the science function visible.** Introduced result-review templates and company-wide showcase sessions with live model demos; migrated model development from EC2/local to AWS SageMaker Studio.

Applied Scientist | Amazon

November 2021 – November 2024 · Berlin, DE

- **Stepped in as interim manager for four months.** With the team left without a manager, took on the role autonomously and ran planning and prioritisation directly with my US-based skip-level on a weekly cadence until the incoming manager arrived.
- **Led two cross-functional, multi-region programmes.** Originated and delivered a Japan-based initiative with a five-person team (2 ML engineers, 1 scientist, 1 TPM, 1 SWE); separately appointed by my manager to own an EU-US programme spanning 10+ software engineers, ML engineers and scientists working within a Tiger Team setting.
- **Owned a production component of Amazon Q and its evaluation framework** — consolidated dataset collection into a single repository with automated quality gates on training and evaluation data, and built a containerised evaluation pipeline delivering reproducible runs with statistical testing.
- **Launched a ranker to production in Amazon Search** with material revenue impact at Amazon Search scale — validated through online A/B testing, with statistically significant relevance improvement (>10 bps).
- **Mentored junior scientists and interns.** An intern gained director-level visibility, published their work at ICML and converted to a full-time Amazon hire; a junior scientist mentee was promoted to mid.
- Trained interviewer; conducted 30+ technical interviews for science positions.

Ph.D. Researcher | Politecnico di Bari

November 2018 – October 2021 · Bari, IT

- Completed doctoral research through October 2021 (awarded cum laude), including scientific consulting on public-private partnership projects delivering research prototypes and formal deliverables.
 - **Junior Applied Scientist, Amazon** (Berlin, DE; July – October 2020) — built a state-of-the-art filter-recommendation solution deployed in production on Amazon.com; published at WWW.

Earlier: Visiting Researcher, Knowledge Media Institute (UK, 2018) · .NET Developer, Fincons Group (IT, 2016–2018).

EDUCATION

- **Ph.D. in Computer Science Engineering** (cum laude), 2018–2021 — Politecnico di Bari, Italy. Dissertation: “Adversarial Machine Learning in Recommender Systems.”
- **M.Sc.** (110 cum laude, 2016–2018) and **B.Sc.** (110 cum laude, 2014–2016) in Computer Engineering — Politecnico di Bari, Italy.

RECOGNITION & SELECTED PUBLICATIONS

- **Industry Chair, RecSys 2026.** 40+ publications across top-tier venues — ICML (1), SIGIR (5), RecSys (3), WWW (2), CIKM (2), WSDM, ECIR, ESWC, KDD — and journals including ACM Computing Surveys and ACM TORS.
- **Awards: Best Short Paper, SIGIR 2025** — “Do LLMs Memorize Recommendation Datasets?” (corresponding author). **Best Paper Award, AdvML @ KDD 2021; Best Short Paper Runner-Up, CIKM 2021.**

TECHNICAL FOUNDATION

- **Data & AI:** LLMs and multi-agent systems, RAG and vector retrieval, ranking and classification, adversarial ML, self-hosted deployment, agentic tooling (Claude Code, OpenAI ADK).
- **Data platform & infrastructure:** Amazon Athena, Iceberg, DynamoDB, SageMaker Studio; data quality and metrics-layer design; Triton serving and containerised evaluation pipelines at tens of millions of requests per day.